

Sistemas Distribuídos

Distributed Transactions

Eurico Pedrosa


António Rui Borges

Universidade de Aveiro - DETI

2025-05-21

Distributed Transactions

What is a Transaction?



- A **transaction** is a **sequence of operations** (reads and writes) issued by a client.
- These operations target **registers in a shared region** managed by a server.
- A transaction must be treated as an **indivisible unit** by the server.



Server Responsibilities

- Ensure **atomic execution** of read/write operations.
- Apply **all-or-nothing semantics**:
 - Save the result as a whole.
 - Or dismiss the entire transaction.
- Handle **failures gracefully**, ensuring either full commit or full rollback.

ACID Properties of Transactions

(Härder and Reuter)

Distributed Transactions



- **Atomicity:** A transaction must be **all or nothing**—either all operations succeed, or none are applied.
- **Consistency:** A transaction must transition the system from **one valid state to another**, preserving invariants.
- **Isolation:** Transactions must be executed **without interference**—their intermediate states must not be visible to other concurrent transactions.
- **Durability:** Once committed, the effects of a transaction must **persist permanently**, even in the event of failures.



- A **transaction** is created and managed by a **coordinator process** on the server side.
- Upon receiving a client request:
 - A transaction is **opened** by the coordinator.
 - A **unique transaction ID** is assigned.
- All subsequent **read-like and write-like operations**:
 - Must include the transaction ID.
 - Return a status:
 - **Success**: The client may continue.
 - **Failure**: The transaction is **aborted**.
- The client may **explicitly abort** the transaction at any time.



- The transaction ends when the client issues an **end-of-transaction** command.
- If all operations succeeded:
 - The transaction is **committed**.
 - All changes become **permanent** in the shared region.
- If any operation failed:
 - An **abort** status is returned.
 - All effects of the transaction are **discarded**.
- After an abort, the client may **retry the transaction** later.



- When **multiple transactions** access or modify the **same registers** in a shared region **concurrently**, **race conditions**- can occur, leading to **data inconsistencies**.
- It is essential to manage concurrent access in a **controlled manner** to ensure correctness.



Conflicting Operations


- Special attention must be given to **pairs of read-like and write-like operations** executed by **different transactions**- on the **same register**.
- If **two transactions attempt to write** to the same register:
 - The register must first be **locked** by one transaction.
 - The modification is made on a **local copy**.
 - When the transaction is **committed**:
 - The shared register is updated.
 - The lock is **released**.
 - Only then can the **second transaction proceed**.



Summary

- Proper synchronization, e.g. **locking mechanisms**, is essential to:
 - Avoid **lost updates** and **inconsistent states**.
 - Ensure **serializability** and **isolation** between transactions.

Distributed Transactions





- In a **distributed transaction**, the **shared region** is partitioned across **multiple servers**.
- Each server manages a portion of the data and **executes part of the transaction**.
- The challenge: How to ensure that **all parts**- of the transaction are **committed together** or **aborted together**?



- A **distributed transaction** can be structured as a **nested transaction**:
 - A **top-level transaction** spawns **subtransactions**.
 - Each subtransaction may **execute on a different machine**.

Key Characteristics

- Subtransactions operate **independently** but are **logically part** of the parent transaction.
- Each subtransaction may **commit locally**, but the effects are **tentative** until the parent commits.



Abort Propagation

- If the **top-level (parent) transaction** aborts:
 - **All subtransactions must also be aborted**, regardless of their local commit status.
 - Ensures **global consistency and atomicity**.

Summary

- Nested transactions provide **modularity** and support for **partial execution** across distributed systems.
- Commit decisions are **hierarchically coordinated** to maintain the **ACID properties**.



To coordinate the outcome across servers, a **two-phase commit protocol**- is used:

Phase 1: Voting

- The **coordinator process** sends a `voteRequest` message to **all participant processes**.
- Each participant responds:
 - `voteCommit` if it is ready to commit its part.
 - `voteAbort` if it cannot commit.



Phase 2: Decision

- The coordinator collects all votes and sends:
 - globalCommit to all participants if **all voted voteCommit**.
 - globalAbort if **at least one voted voteAbort**.
- Participation Action Upon receiving the final decision:
 - If globalCommit: each participant **commits** its local part.
 - If globalAbort: each participant **aborts** and discards its local changes.


Two-Phase Commit Protocol (2PC)

(without failures)

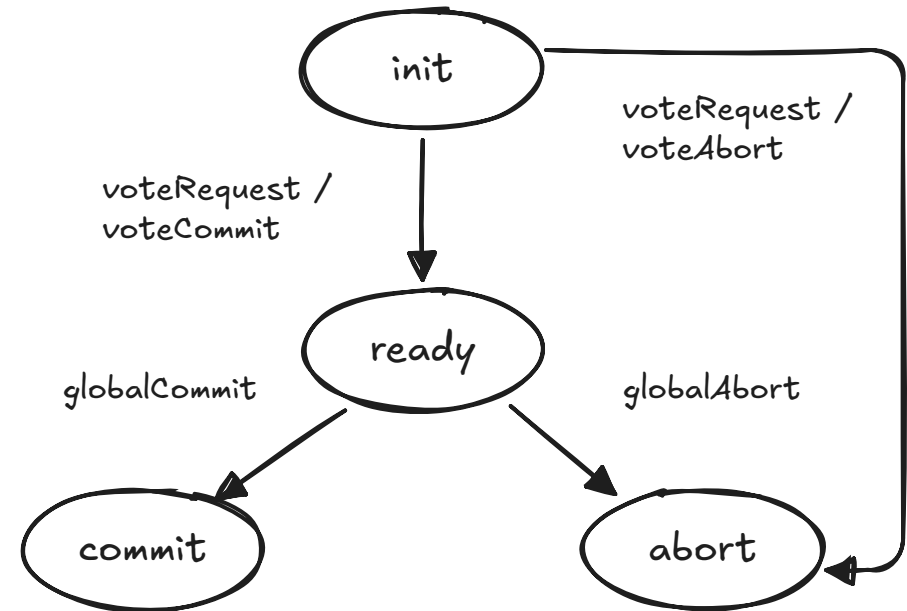
Distributed Transactions



Coordinator Process



Participant Process




Two-Phase Commit Protocol (2PC)

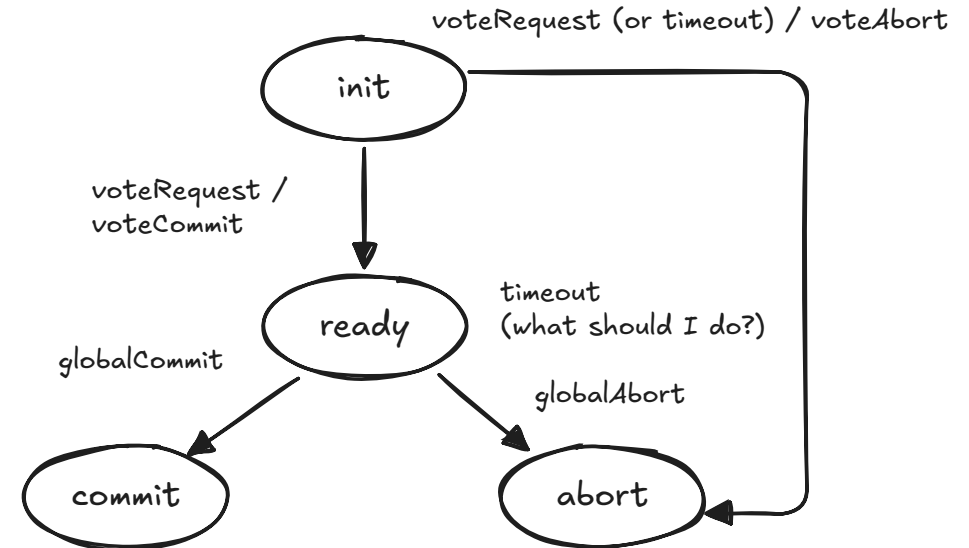
(with failures)



Coordinator Process



Participant Process





- **3PC** extends the **Two-Phase Commit (2PC)** protocol to avoid **blocking** in the event of a coordinator crash.
- Introduces an **intermediate phase** to ensure that participants can make progress without indefinite waiting.

Phase 1: Voting

- The **coordinator** sends a `voteRequest` to all **participants**.
- Each participant replies with:
 - `voteCommit` (ready to commit), or
 - `voteAbort` (cannot commit).



Phase 2: Pre-Commit

- If all participants vote `voteCommit`, the coordinator sends a `preCommit` message.
- Participants:
 - Acknowledge readiness.
 - Enter a **prepared state** (promise to commit but do not commit yet).

Phase 3: Do-Commit

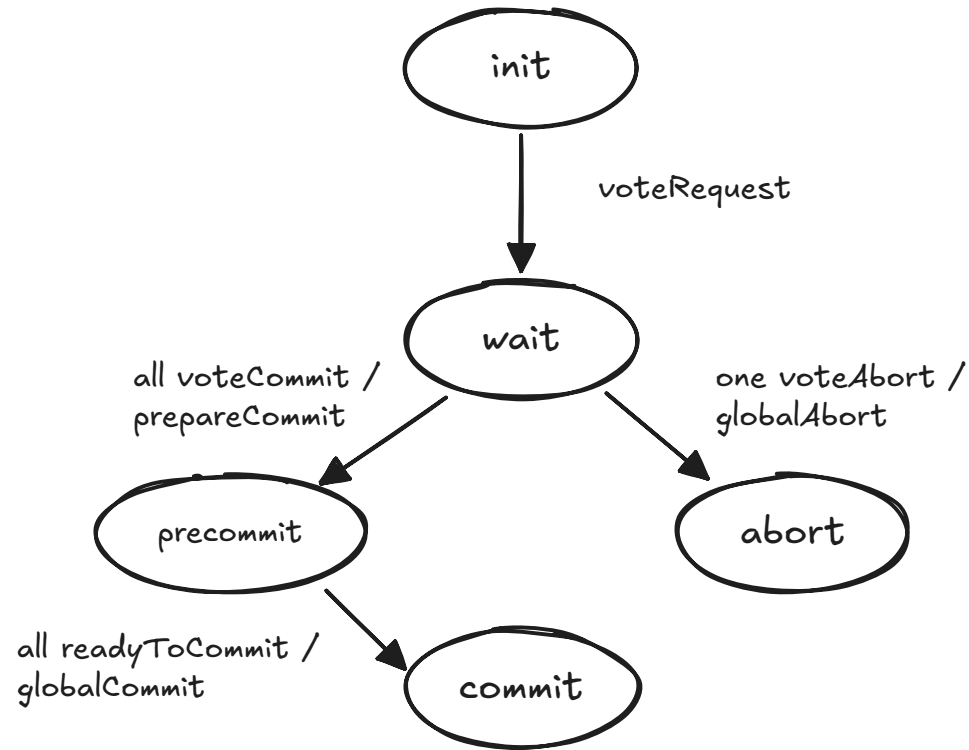
- After receiving all acknowledgments, the coordinator sends a `globalCommit`.
- Participants **perform the actual commit**.

Three-Phase Commit (3PC)

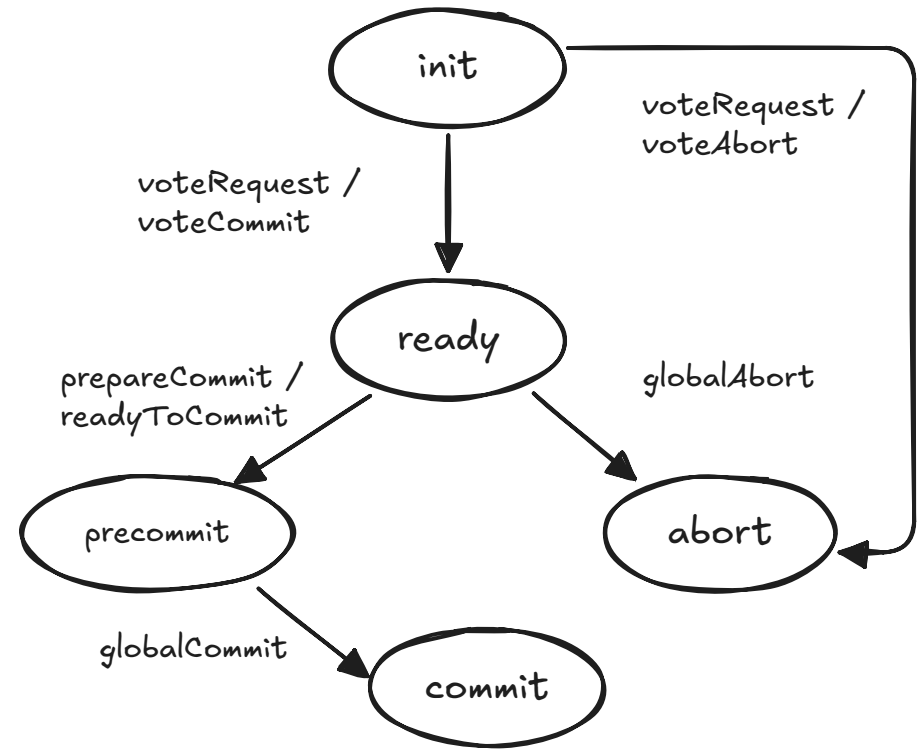
(without failures)



Coordinator Process



Participant Process



Three-Phase Commit (3PC)

Safety and Non-Blocking Guarantees

Distributed Transactions



- **Safe intermediate states** prevent direct transitions to COMMIT or ABORT.
- **Progress is always possible**: no state requires indefinite waiting on others.

Failure Handling

- If the **coordinator crashes during PreCommit**:
 - Participants can safely **commit** (commit is known to be agreed upon).
- If a **participant times out**, it checks others:



- ▶ If others are in COMMIT or ABORT: **adopt that decision.**
- ▶ If all are in PRECOMMIT: **safely commit.**
- ▶ If all are in READY: **abort to ensure safety.**

Why 3PC is Rarely Used

- Adds **significant complexity and overhead.**
- In practice, **2PC suffices** for most use cases in well-managed and reliable systems.

Suggested Reading



- M. van Steen and A.S. Tanenbaum, Distributed Systems, 4th ed., distributed-systems.net, 2023.