

Sistemas Distribuídos

Summary

Eurico Pedrosa

António Rui Borges

Universidade de Aveiro - DETI

2025-05-29

Introduction to Distributed Systems



- Understand principles and practical design of distributed systems
 - Communication and synchronization
 - Java-based implementations
- Develop skills to build distributed applications



- Grasp design issues in distributed systems
- Tackle concurrency and consistency in real systems
- Apply Java RMI and message passing

Concurrency

- Computer Architecture
 - CPU
 - Main Memory
 - Pipelining
- Program vs Process
 - Process State
- Multithreading
 - Organization

System Models

- **Client-Server:** centralized resource control
- **Peer-to-Peer:** equal roles, replication
- **Publisher-Subscriber:** broker decouples producers and consumers

- **Interaction:** bandwidth, latency, jitter
- **Failure:** omission, timing, byzantine
- **Security:** process and channel threats, access control

Message Passing and Communication

Communication Fundamentals



- **Latency, transfer rate, bandwidth**
- **Synchronous** vs **asynchronous** primitives
- **Blocking** vs **non-blocking** operations



- TCP: reliable, bidirectional, connection-based
- UDP: connectionless, low-overhead
- Socket identified by IP and port

Distributed Execution



- Marshaling/unmarshaling of data
- Stub acts as proxy for remote objects
- Communication via structured messages



- Server base thread + proxy agent
- Mixed architecture: server also acts as client
- Java serialization simplifies implementation

Concurrency in Distributed Systems



- **Independent** vs **cooperating** processes
- **Critical regions**: mutual exclusion essential
- **Deadlock** and **indefinite postponement**



- **Monitors:** encapsulated access with wait/signal
- **Semaphores:** general-purpose mutual exclusion
- Java concurrency tools: barriers, locks, atomic ops

Synchronization and Time



- **Global** vs **local** vs **logical** time
- Synchronization limits due to drift and latency



- Cristian's Method (UTC server)
- Berkeley Algorithm (internal synchronization)
- NTP: hierarchical, resilient synchronization over Internet

Logical Clocks

- Capture **happened-before** relation
- Scalar timestamps ensure partial ordering
- Extended timestamps for total ordering

- Precise causality tracking
- Each process maintains vector of logical times
- Allows detecting concurrency vs causality

Group Communication



- **Centralized:** guardian manages access
- **Token Ring:** token grants access in sequence
- **Ricart & Agrawala:** logical clocks for total order
- **Maekawa:** voting subsets to minimize messages



- **Ring-based** and **unstructured group** algorithms
- Key properties: termination, unambiguity, consensus
- Failure handling: timeout, retries, ACK-based recovery

Consistency and Replication



- **Strict Consistency:** ideal but unrealistic
- **Linearizability:** respects real-time order
- **Sequential Consistency:** respects program order
- **Causal Consistency:** enforces causal relations
- **FIFO Consistency:** per-process order preserved



- Use of logical clocks (scalar or vector)
- Propagation of write-like operations
- Application-dependent trade-offs

Distributed Transactions



- **Atomicity, Consistency, Isolation, Durability**
- Local commit or global abort
- Locking to prevent race conditions



- **2-Phase Commit (2PC):**
 - Voting → Decision (commit/abort)
- **3-Phase Commit (3PC):**
 - Adds Pre-commit phase
 - Ensures progress and non-blocking behavior



- Hierarchical coordination of sub-transactions
- Abort propagates to all sublevels